

Stop prompting. Start orchestrating.

The step that turns Claude Code from a fast assistant into an engineering process.

solo.

One context is one opinion

A single agent writes the code, then grades its own work. It defends its first idea to the end of the context window — the same trap that makes authors bad reviewers of their own code.

What the orchestration step looks like

SCOUT	→ map the codebase in parallel
FAN OUT	→ one agent per concern
VERIFY	→ skeptics refute each finding
MERGE	→ one ranked synthesis

When to reach for it

- 01 The work outgrows one context
- 02 Findings need independent checks
- 03 The task splits along clean seams
- 04 Being wrong is expensive

Orchestrate

Codebase audits

Wide migrations

Adversarial review

Research sweeps


Stay solo

One-file fixes

Quick questions

Exploration


Anything cheap to redo



Orchestration isn't speed. It's doubt.

Parallel agents don't just finish sooner. They disagree
— and the disagreement is the signal.

“An agent reviewing its own work is one opinion counted twice.”



Ship the process, not the prompt.

Want the full Night Ace orchestra for Claude Code — agents, commands, and the playbook? Comment “ACE” and we'll send you the package.